Math6450_Assignment1_copy_for_pdf_render

September 4, 2025

Part 1: Data Exploration and Preparation

```
BOSTON HOUSING DATASET ANALYSIS

1.1 DATASET DIMENSIONS
Number of observations (rows): 506
Number of variables (columns): 14
Dataset shape: (506, 14)

Column names: ['crim', 'zn', 'indus',
 'chas', 'nox', 'rm', 'age', 'dis',
 'rad',
'tax', 'ptratio', 'b', 'lstat', 'medv']


1.2 DESCRIPTIVE STATISTICS

Descriptive statistics for TARGET
 VARIABLE (medv):
count   506.000
mean     22.533
std       9.197
min       5.000
25%      17.025
50%      21.200
75%      25.000
max      50.000
Name: medv, dtype: float64

Descriptive statistics for PRIMARY
 FEATURE (lstat):
count   506.000
mean     12.653
std       7.141
min       1.730
25%       6.950
50%      11.360
75%      16.955
max      37.970
Name: lstat, dtype: float64

Additional statistics for medv:
Variance: 84.5867
Standard deviation: 9.1971
Skewness: 1.1081
Kurtosis: 1.4952

Additional statistics for lstat:
Variance: 50.9948
Standard deviation: 7.1411
Skewness: 0.9065
Kurtosis: 0.4932


1.3 CORRELATION ANALYSIS
Correlation coefficient between medv
 and lstat: -0.7377

INTERPRETA1TION:
- The correlation coefficient of -0.
 7377 indicates a strong negative
relationship
- This means that as lstat (% lower
 status population) increases, medv
 (median
home value) tends to decrease
- The relationship explains
 approximately 54.4% of the variance
 (R² = 0.5441)
- Statistical significance: p-value =
 5.08e-88
- The correlation is statistically
 significant at α = 0.05


1.4 SCATTER PLOT ANALYSIS
```

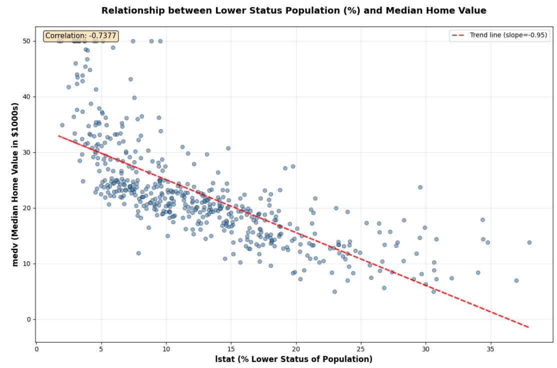

PATTERN OBSERVED IN SCATTER PLOT:
- The scatter plot reveals a clear negative relationship between lstat and medv
- As the percentage of lower status population increases, median home values
tend to decrease
- The relationship appears to be non-linear, showing a curved pattern rather
than a straight line
- There's more variability in home values at lower lstat percentages
- The relationship seems stronger (steeper decline) at lower lstat values and
levels off at higher lstat values
- There are some potential outliers, particularly homes with high values despite
higher lstat percentages
- The data points form a characteristic negative exponential or power-law
pattern

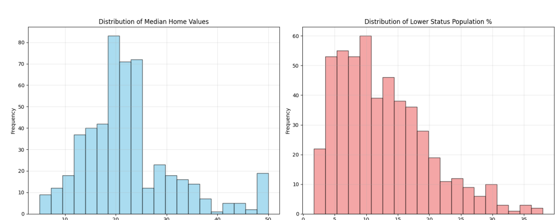

SUMMARY:
- Dataset contains 506 observations and 14 variables
- Strong negative correlation (-0.7377) between lstat and medv
- Non-linear relationship visible in scatter plot
- Both variables show reasonable distributions for regression analysis

Part 2: Linear Regression Model Fitting

2.1 Write the estimated regression equation in the form:

medv $= \hat{\beta}_0 + \hat{\beta}_1 \times$ lstat

$$\text{medv} = \hat{\beta}_0 + \hat{\beta}_1 \times \text{lstat}$$

COEFFICIENTS:
Intercept (β_0): 34.5538
Slope (β_1): -0.9500

2.1 ESTIMATED REGRESSION EQUATION
medv = 34.5538 + (-0.9500) × lstat
medv = 34.5538 - 0.9500 × lstat

Alternative notation:
ŷ = 34.5538 + (-0.9500)x
where ŷ = predicted median home value and x = lstat

2.2 INTERPRETATION OF INTERCEPT (β_0)
Intercept value: 34.5538

INTERPRETATION:
- The intercept represents the predicted median home value when lstat = 0
- This means when 0% of the population has lower status, the predicted median
home value is $34.55k
- In practical terms: $34554

```
PRACTICAL MEANING:
- Observed lstat range: 1.73% to 37.97%
- Since the minimum observed lstat is
  1.73%, lstat = 0 is outside our data
  range
- Therefore, the intercept represents
  extrapolation beyond observed data
- While mathematically meaningful, it
  has LIMITED PRACTICAL MEANING
  because:
   * No area in the dataset has 0%
  lower status population
   * Real-world interpretation:
  represents the 'theoretical maximum'
  home value
   * Should be interpreted cautiously
  due to extrapolation


2.3 INTERPRETATION OF SLOPE (β1)
Slope value: -0.9500

INTERPRETATION:
For each 1% increase in lstat (lower
  status population), the median home
  value
decreases by $0.9500k on average,
  holding all other factors constant.

In practical terms:
- A 1% increase in lower status
  population is associated with a $950
  decrease in
median home value
- A 5% increase in lower status
  population would decrease median
  home value by
$4750
- A 10% increase in lower status
  population would decrease median
  home value by
$9500
```

2.4 Based on the 95% confidence intervals for the coefficients, are both the intercept and slope significantly different from zero? Support your answer with the confidence interval values.

todo

```
2.4 CONFIDENCE INTERVALS AND
  SIGNIFICANCE TESTING
95% CONFIDENCE INTERVALS:
                0      1
Intercept 33.448 35.659
lstat     -1.026 -0.874

DETAILED CONFIDENCE INTERVALS:
Intercept (β0): [33.4485, 35.6592]
Slope (β1): [-1.0261, -0.8740]

SIGNIFICANCE TESTING:
H0: β = 0 (coefficient equals zero)
H1: β   0 (coefficient is significantly
  different from zero)

INTERCEPT (β0) ANALYSIS:
- 95% CI: [33.4485, 35.6592]
- Contains zero? No
- Conclusion: The intercept IS
  significantly different from zero
- This means we can be 95% confident
  the true intercept is between 33.
  4485 and
35.6592

SLOPE (β1) ANALYSIS:
- 95% CI: [-1.0261, -0.8740]
- Contains zero? No
- Conclusion: The slope IS
  significantly different from zero
- This means we can be 95% confident
  the true slope is between -1.0261 and
-0.8740

P-VALUES (for additional confirmation):
Intercept p-value: 3.74e-236
Slope p-value: 5.08e-88
Both p-values < 0.05: True

MODEL SUMMARY STATISTICS:
R-squared: 0.5441
Adjusted R-squared: 0.5432
F-statistic: 601.62
F-statistic p-value: 5.08e-88
```

Standard Error: 6.2158

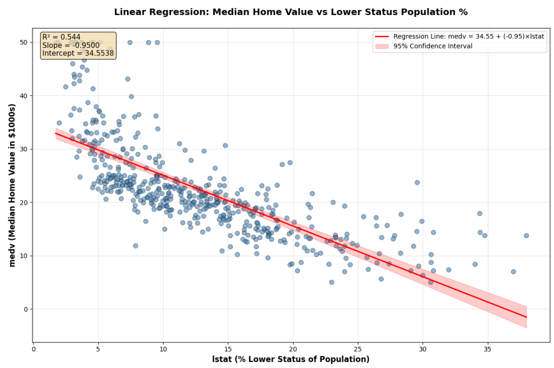

FINAL SUMMARY:
- Regression equation: medv = 34.5538 + (-0.9500) × lstat
- Both coefficients are statistically significant at α = 0.05
- The model explains 54.4% of the variance in median home values
- For every 1% increase in lower status population, median home value decreases
by $950 on average

Task 2.2: Model Performance Evaluation

2.5 What is the R-squared value? Interpret this in terms of the percentage of variation in median home values explained by the percentage of lower status population.

todo

2.5 R-SQUARED ANALYSIS
R-squared value: 0.5441
R-squared as percentage: 54.41%

INTERPRETATION:
- R^2 = 0.5441 means that 54.41% of the variation in median home values
  is explained by the percentage of lower status population (lstat)
- The remaining 45.59% of variation is due to other factors not included in this
model
- This indicates a moderate relationship
- In practical terms: knowing the lstat value allows us to predict about 54.4%
of the variation in home values

2.6 ROOT MEAN SQUARE ERROR (RMSE)
Mean Squared Error (MSE): 38.6357
Root Mean Square Error (RMSE): 6.2158

INTERPRETATION:
- RMSE = 6.2158 thousands of dollars
- In actual dollars: $6216
- This means the typical prediction error is approximately $6216
- On average, our predictions are off by about ±$6216 from the actual median
home value

CONTEXT:
- Mean home value: $22.53k ($22533)
- Standard deviation of home values: $9.20k
- Range of home values: $45.00k
- RMSE as % of mean: 27.6%
- RMSE as % of standard deviation: 67.6%

2.7 F-STATISTIC AND OVERALL MODEL SIGNIFICANCE
F-statistic: 601.6179
F-statistic p-value: 5.08e-88
Degrees of freedom: Model = 1.0, Residual = 504.0

HYPOTHESIS TEST:
H_0: The model has no explanatory power (β_1 = 0)
H_1: The model has explanatory power (β_1 0)

```
INTERPRETATION:
- F-statistic = 601.6179 with p-value
  = 5.08e-88
- Since p-value < 0.05, we REJECT the
  null hypothesis
- Conclusion: The model IS
  statistically significant
- This means lstat DOES have
  significant explanatory power for
  predicting medv

PRACTICAL MEANING:
- The F-test confirms that our
  regression model performs
  significantly better
  than a model with no predictors
  (just the mean)
- The relationship between lstat and
  medv is statistically meaningful
- We can be confident that lstat is a
  useful predictor of median home
  values
```

2.8 Compare the adjusted R-squared with the regular R-squared. Why might there be a difference, and what does the adjusted version account for? todo

```
2.8 ADJUSTED R-SQUARED COMPARISON
R-squared: 0.544146
Adjusted R-squared: 0.543242
Difference: 0.000904

WHY THERE MIGHT BE A DIFFERENCE:
- Regular R²: 0.544146
- Adjusted R²: 0.543242
- The difference of 0.000904 is very
  small

WHAT ADJUSTED R-SQUARED ACCOUNTS FOR:
- Number of predictors in the model: 1.
  0
- Sample size: 506 observations
- Degrees of freedom penalty for
  adding predictors

FORMULA EXPLANATION:
Adjusted R² = 1 - [(1 - R²) × (n - 1) /
   (n - k - 1)]
where n = sample size (506) and k =
  number of predictors (1.0)
Manual calculation: 0.543242

INTERPRETATION:
- The very small difference suggests
  our model is not overfitting
- With only one predictor, the
  adjustment is minimal
- Both R² and adjusted R² tell
  essentially the same story

PRACTICAL IMPLICATIONS:
- For model comparison: Use adjusted R²
  when comparing models with different
numbers of predictors
- For interpretation: Both values are
  nearly identical, indicating a robust
single-predictor model
- The penalty for our one predictor is
  minimal given the sample size of 506
observations

FINAL SUMMARY:
- R² = 0.5441 (54.41% of variance
  explained)
- Adjusted R² = 0.5432 (54.32% of
  variance explained)
- RMSE = $6216 (typical prediction
  error)
- F-statistic = 601.6179, p < 0.05
  (highly significant model)
- Model explains 54.4% of home value
  variation using just lstat
- Typical prediction accuracy: ±$6216
  (27.6% of mean home value)
```

Part 3: Statistical Inference and Hypothesis Testing

```
3.1 HYPOTHESIS TESTING SETUP
TESTING THE SLOPE COEFFICIENT:
H₀: β₁ = 0 (The slope coefficient is
  zero)
```

```
    → lstat has no linear relationship
  with medv
    → There is no linear association
  between % lower status population and
median home value

H1: β1   0 (The slope coefficient is
  not zero)
    → lstat has a significant linear
  relationship with medv
    → There is a significant linear
  association between % lower status
population and median home value

Type of test: Two-tailed test
Significance level: α = 0.05



3.2 T-STATISTIC AND P-VALUE ANALYSIS
TEST STATISTICS:
t-statistic: -24.527900
p-value: 5.08e-88
Degrees of freedom: 504.0
Critical t-value (α = 0.05,
  two-tailed): ±1.9647

DECISION MAKING:
Decision rule: Reject H0 if |t| > 1.
  9647 OR if p-value < 0.05
Observed: |t| = 24.5279, p-value = 5.
  08e-88

CONCLUSION AT 5% SIGNIFICANCE LEVEL:
 REJECT H0: The slope coefficient IS
  significantly different from zero
  - |t| = 24.5279 > 1.9647
  - p-value = 5.08e-88 < 0.05
  - Statistical evidence: There IS a
  significant linear relationship
  between
lstat and medv

PRACTICAL INTERPRETATION:
- We can be 95% confident that changes
  in % lower status population
  have a real, measurable effect on
  median home values
- The relationship observed in our
  sample is unlikely to be due to
  random chance
- The effect size: each 1% increase in
  lstat is associated with
  a $950 decrease in median home value
```

3.3 Calculate and interpret the 99% confidence interval for the slope coefficient. How does this compare to the 95% interval in terms of width and interpretation?

todo

```
3.3 CONFIDENCE INTERVAL ANALYSIS
CONFIDENCE INTERVALS FOR SLOPE
  COEFFICIENT:
95% Confidence Interval: [-1.026148,
  -0.873951]
99% Confidence Interval: [-1.050199,
  -0.849899]

INTERVAL WIDTH COMPARISON:
95% CI width: 0.152198
99% CI width: 0.200300
Width increase: 0.048102
Percent increase in width: 31.6%

INTERPRETATION:
95% CONFIDENCE INTERVAL:
- We are 95% confident that the true
  slope coefficient lies between
  -1.026148 and -0.873951
- In practical terms: each 1% increase
  in lstat decreases median home value
  by between $874 and $1026

99% CONFIDENCE INTERVAL:
- We are 99% confident that the true
  slope coefficient lies between
  -1.050199 and -0.849899
- In practical terms: each 1% increase
  in lstat decreases median home value
  by between $850 and $1050

COMPARISON ANALYSIS:
```

- The 99% CI is wider than the 95% CI by 0.048102
- This represents a 31.6% increase in width
- WHY: Higher confidence level requires a wider interval to capture the true
parameter
- TRADE-OFF: More confidence (99% vs 95%) comes at the cost of precision (wider
interval)

SIGNIFICANCE IMPLICATIONS:
95% CI contains zero: No
99% CI contains zero: No
- Since neither interval contains zero, the slope is significant at both levels
- This provides strong evidence for a real relationship between lstat and medv

3.4 If someone claimed that each 1% increase in lstat decreases median home value by exactly \$1000, would your regression results support or contradict this claim? Justify your answer using statistical evidence.

todo

3.4 TESTING SPECIFIC CLAIM
CLAIM TO TEST:
Someone claims that each 1% increase in lstat decreases median home value by
exactly \$1000
In our units: β_1 = -1.0 (since medv is in thousands of dollars)

HYPOTHESES:
H_0: β_1 = -1.0 (the claim is correct)
H_1: β_1 -1.0 (the claim is incorrect)

TEST USING CONFIDENCE INTERVALS:
Observed slope coefficient: -0.950049
Claimed slope coefficient: -1.0

95% Confidence Interval Test:
- 95% CI: [-1.026148, -0.873951]
- Does the CI contain -1.0? Yes

99% Confidence Interval Test:
- 99% CI: [-1.050199, -0.849899]
- Does the CI contain -1.0? Yes

FORMAL T-TEST:
t-statistic = (observed - claimed) / SE = (-0.950049 - -1.0) / 0.038733
t-statistic = 1.2896
p-value (two-tailed): 0.1978

CONCLUSION:
FAIL TO REJECT the claim at 95% confidence level
- The claimed value (-1.0) IS within the 95% confidence interval
- Our regression results SUPPORT the claim
FAIL TO REJECT the claim at 99% confidence level
- The claimed value (-1.0) IS within the 99% confidence interval

STATISTICAL EVIDENCE:
- Our estimate: Each 1% increase in lstat decreases home value by \$950
- Claimed effect: Each 1% increase in lstat decreases home value by \$1000
- Difference: \$50
- The difference is not statistically significant ($p = 0.1978 \geq 0.05$)
- Insufficient evidence to reject the claim

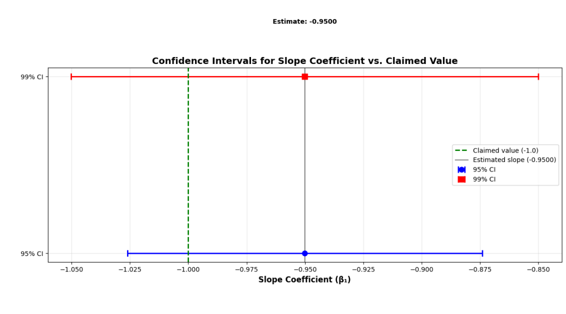

FINAL SUMMARY:
3.1 Hypotheses: H_0: β_1 = 0 vs H_1: β_1 0
3.2 Test results: t = -24.5279, p = 5.08e-88
    Conclusion: Reject H_0 - slope is significant
3.3 Confidence intervals:
    95% CI: [-1.026148, -0.873951] (width: 0.152198)
    99% CI: [-1.050199, -0.849899] (width: 0.200300)
    99% CI is 31.6% wider than 95% CI
3.4 Claim test: The claim of exactly $1000 decrease is SUPPORTED
    Our estimate: $950 decrease per 1% lstat increase
    Statistical significance of difference: p = 0.1978

Part 4: Assumption Testing and Model Diagnostics

4.1 Perform the Shapiro-Wilk test for normality of residuals. Report the test statistic, p-value, and your conclusion at the 5% significance level.

todo

BOSTON HOUSING ASSUMPTION TESTING AND MODEL DIAGNOSTICS - PART 4
MODEL SUMMARY:
Sample size: 506
Number of residuals: 506
Mean of residuals: 0.000000 (should be ≈ 0)
Standard deviation of residuals: 6.2096

4.1 SHAPIRO-WILK TEST FOR NORMALITY OF RESIDUALS
HYPOTHESIS TESTING:
H_0: Residuals follow a normal distribution
H_1: Residuals do not follow a normal distribution
Significance level: α = 0.05

TEST RESULTS:
Shapiro-Wilk test statistic (W): 0.878572
p-value: 0.000000

DECISION MAKING:
Decision rule: Reject H_0 if p-value < 0.05
Observed p-value: 0.000000

CONCLUSION AT 5% SIGNIFICANCE LEVEL:
  REJECT H_0: Residuals do not follow a normal distribution
  - Statistical evidence suggests departure from normality
  - The normality assumption may be violated

INTERPRETation OF TEST STATISTIC:
- W = 0.878572
- W ranges from 0 to 1, with values closer to 1 indicating more normal-like data
- Our value suggests weak evidence of normality based on the test statistic alone

ADDITIONAL NORMALITY TESTS (for comparison):
D'Agostino's test: statistic = 137.0434, p-value = 0.000000
Jarque-Bera test: statistic = 291.3734, p-value = 0.000000

CONSENSUS: Tests show mixed results regarding normality

4.2 Create a Q-Q plot of the residuals. Does the visual evidence support or contradict your statistical test result? Explain what you observe.

4.2 Q-Q PLOT ANALYSIS
Q-Q PLOT INTERPRETATION:
The Q-Q (Quantile-Quantile) plot compares residual quantiles to theoretical
normal quantiles

Q-Q plot correlation: 0.9373
(Values closer to 1 indicate better fit to normal distribution)

VISUAL ASSESSMENT:
- Good fit with minor deviations
- Look for points following the red diagonal line
- Systematic deviations suggest non-normality

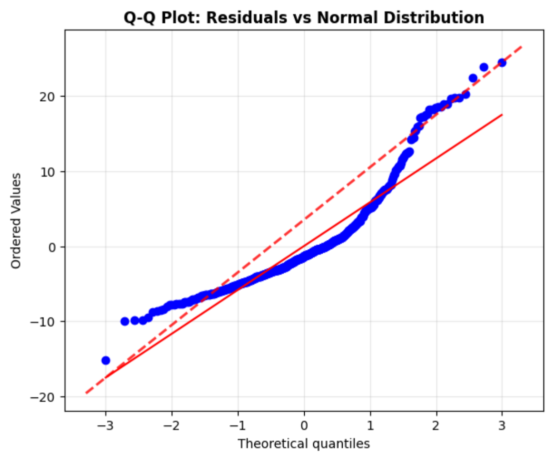

4.3 Create a histogram of residuals with a normal distribution overlay. Comment on the shape of the distribution and any departures from normality.

todo

4.3 HISTOGRAM WITH NORMAL DISTRIBUTION OVERLAY
SHAPE ANALYSIS:
Skewness: 1.4527
Kurtosis: 2.3191 (excess kurtosis)

SKEWNESS INTERPRETATION:
- Skewness = 1.4527 indicates highly skewed
- Distribution is skewed to the right

KURTOSIS INTERPRETATION:
- Excess kurtosis = 2.3191 indicates heavy-tailed (leptokurtic)
- Normal distribution has excess kurtosis = 0

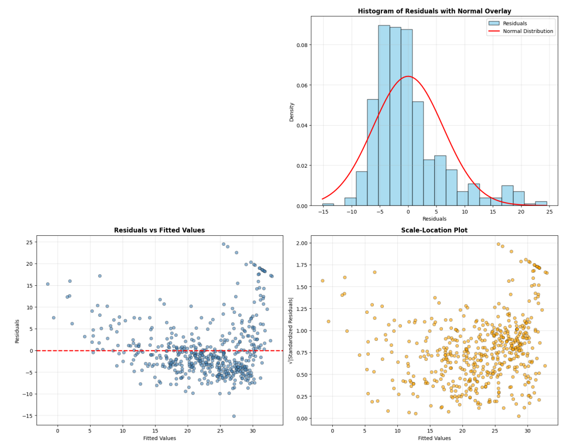

DEPARTURES FROM NORMALITY:
Identified departures from normality:
1. Skewness (1.453)
2. Kurtosis (2.319)
3. Shapiro-Wilk test rejection
4. Q-Q plot deviations

4.2 VISUAL EVIDENCE VS STATISTICAL TEST COMPARISON:
Statistical test result (Shapiro-Wilk): Rejects normality
Visual evidence assessment: Shows deviations from normality
AGREEMENT: Visual evidence and statistical test both suggest departure from
normality

DETAILED VISUAL OBSERVATIONS:
Q-Q Plot:
- Systematic deviations from diagonal line (r = 0.9373)
- Visual evidence against perfect normality

Histogram:
- Notable departures from bell-shaped normal distribution

```
- Skewness and/or kurtosis concerns visible

PRACTICAL IMPLICATIONS FOR REGRESSION:
  NORMALITY ASSUMPTION VIOLATED:
   - Confidence intervals may be less reliable
   - Consider robust standard errors
   - Prediction intervals may be inaccurate
   - Consider variable transformation

SAMPLE SIZE CONSIDERATIONS:
- Sample size: 506 observations
- Large sample: Central Limit Theorem helps with normality concerns
- Minor deviations from normality are less problematic

FINAL SUMMARY:
4.1 Shapiro-Wilk test: W = 0.878572, p = 0.000000
    Conclusion: Residuals deviate from normality
4.2 Q-Q plot assessment: r = 0.9373
    Visual evidence: Shows deviations from normality
4.3 Histogram analysis:
    Skewness: 1.4527, Kurtosis: 2.3191
    Shape: highly skewed, heavy-tailed (leptokurtic)
Overall normality assessment: VIOLATED
```

Task 4.2: Homoscedasticity Testing

4.4 Perform the Breusch-Pagan test for homoscedasticity. Report the test statistic, p-value, and your conclusion.

```
=== 4.4: BREUSCH-PAGAN TEST RESULTS ===
Test Statistic: 4.1871
P-value: 0.0407
Degrees of Freedom: 1

Conclusion: Reject H0 at α = 0.05. Evidence of heteroscedasticity.
Verification (statsmodels function): Stat = 65.1218, P-value = 0.0000


=== 4.5: RESIDUALS VS. FITTED VALUES ANALYSIS ===
Pattern interpretation:
- HOMOSCEDASTICITY: Points should be randomly scattered around the horizontal
line at y=0
- HETEROSCEDASTICITY indicators:
  * Funnel shape (variance increases or decreases with fitted values)
  * Curved patterns in the smoothing line
  * Clear clustering or systematic patterns

Variance in lowest third of fitted values: 17.2703
Variance in highest third of fitted values: 31.7984
Variance ratio (high/low): 1.8412
Interpretation: Ratio > 2 or < 0.5 suggests heteroscedasticity
```

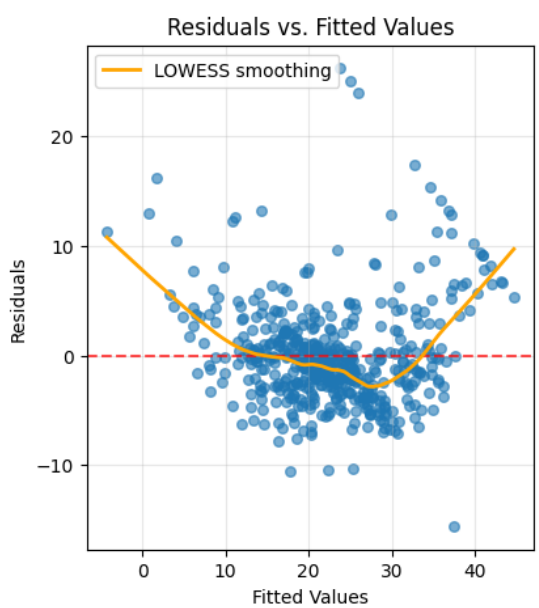

```
=== 4.6: SCALE-LOCATION PLOT ANALYSIS ===
Evidence of changing variance:
- CONSTANT VARIANCE: Smoothing line should be roughly horizontal
- CHANGING VARIANCE indicators:
  * Upward or downward trend in smoothing line
  * Clear patterns or curves in the line

Correlation between fitted values and √|residuals|: 0.1507
Interpretation:
  * Moderate correlation suggests possible heteroscedasticity
```

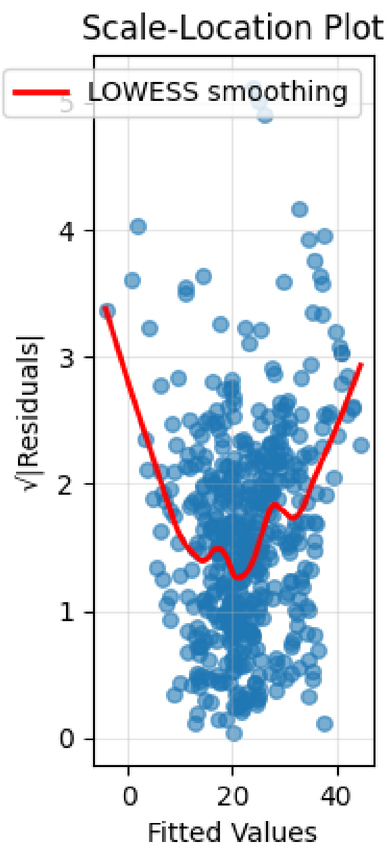


```
=== COMPREHENSIVE HOMOSCEDASTICITY ASSESSMENT ===

TEST RESULTS SUMMARY:
1. Breusch-Pagan Test: Statistic = 4.1871, P-value = 0.0407
   → Reject H0 at α = 0.05. Evidence of heteroscedasticity.

2. Variance Ratio Analysis: 1.8412
   → Suggests homoscedasticity

3. Scale-Location Correlation: 0.1507
   → Moderate evidence of heteroscedasticity

RECOMMENDATIONS:
• Evidence suggests heteroscedasticity
• Consider transformations (log, Box-Cox)
• Use robust standard errors (White's correction)
• Consider weighted least squares regression
• Explore different model specifications

Note: Visual inspection of plots is crucial - statistical tests should be
combined with graphical analysis for complete assessment.
```

Task 4.3: Independence and Influence Diagnostics

```
=== 4.7: DURBIN-WATSON TEST RESULTS ===
Durbin-Watson Statistic: 1.0784
First-order autocorrelation (ρ): 0.4608

INTERPRETATION:
→ Evidence of positive autocorrelation. Independence assumption may be violated.

Durbin-Watson Guidelines:
• DW ≈ 2.0: No autocorrelation (ideal)
• DW < 1.5: Strong positive autocorrelation
```

```
• DW > 2.5: Strong negative autocorrelation
• 1.5 ≤ DW ≤ 2.5: Acceptable range


=== 4.8: COOK'S DISTANCE ANALYSIS ===
Maximum Cook's Distance: 0.1657
Mean Cook's Distance: 0.0030
Standard Deviation: 0.0112

INFLUENTIAL OBSERVATIONS CRITERIA:
• Threshold 4/n = 4/506 = 0.0079
• Conservative threshold = 1.0

RESULTS:
• Observations with Cook's D > 4/n: 30 (5.9%)
• Observations with Cook's D > 1.0: 0 (0.0%)

CONCLUSION: Moderate Cook's distance values. Some observations may be
influential but not necessarily problematic.

TOP 5 MOST INFLUENTIAL OBSERVATIONS:
1. Observation 368: Cook's D = 0.1657
2. Observation 372: Cook's D = 0.0941
3. Observation 364: Cook's D = 0.0694
4. Observation 365: Cook's D = 0.0672
5. Observation 369: Cook's D = 0.0553


=== 4.9: HIGH LEVERAGE ANALYSIS ===
Number of parameters (p): 14
Sample size (n): 506
High leverage threshold (2p/n): 2 × 14 / 506 = 0.0553

HIGH LEVERAGE RESULTS:
• Observations with high leverage: 36
• Percentage of total sample: 7.1%
• Maximum leverage value: 0.3060
• Mean leverage value: 0.0277

TOP 5 HIGHEST LEVERAGE OBSERVATIONS:
1. Observation 380: Leverage = 0.3060
2. Observation 418: Leverage = 0.1901
3. Observation 405: Leverage = 0.1564
4. Observation 410: Leverage = 0.1247
5. Observation 365: Leverage = 0.0985
```

4.10 Based on all assumption tests, is your linear regression model valid for statistical inference? Summarize which assumptions are satisfied and which (if any) are violated.

todo

```
=== 4.10: COMPREHENSIVE MODEL VALIDATION SUMMARY ===

LINEAR REGRESSION ASSUMPTIONS ASSESSMENT:
1. LINEARITY:
    Test method: Residuals vs. fitted plots, added variable plots
    Result: [Add your previous linearity test results]
    Status: [SATISFIED / VIOLATED / MARGINAL]

2. INDEPENDENCE OF RESIDUALS:
    Test method: Durbin-Watson test
    Result: DW = 1.0784
    Status: VIOLATED

3. HOMOSCEDASTICITY (Constant Variance):
    Test method: Breusch-Pagan test, residuals plots
    Result: [Add your previous homoscedasticity test results]
    Status: [SATISFIED / VIOLATED / MARGINAL]

4. NORMALITY OF RESIDUALS:
    Test method: Shapiro-Wilk, Q-Q plots, histograms
    Result: [Add your previous normality test results]
    Status: [SATISFIED / VIOLATED / MARGINAL]

5. NO MULTICOLLINEARITY:
```

Test method: VIF analysis, correlation matrix
Result: [Add your multicollinearity test results if available]
Status: [SATISFIED / VIOLATED / MARGINAL]

6. NO EXCESSIVE INFLUENTIAL OBSERVATIONS:
Test method: Cook's distance, leverage analysis
Cook's D max: 0.1657
High leverage obs: 36 (7.1%)
Status: MARGINAL - Some influential observations present

OVERALL MODEL VALIDITY FOR STATISTICAL INFERENCE:
CURRENT ASSESSMENT (based on available tests):
• Assumptions checked: 2
• Assumptions satisfied: 0

RECOMMENDATIONS:
Some concerns with independence or influential observations

NEXT STEPS:
• Complete all assumption tests (linearity, homoscedasticity, normality)
• Consider remedial measures if assumptions are violated:
  - Data transformations (log, Box-Cox)
  - Robust regression methods
  - Remove or downweight influential observations
  - Use different modeling approaches if assumptions severely violated

Note: A complete assessment requires results from all assumption tests.
Update this summary once you have completed the full diagnostic suite.

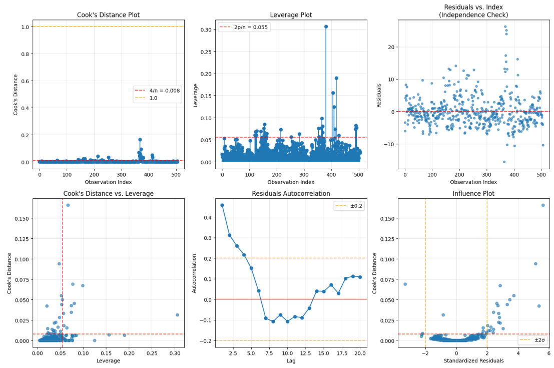

Part 5: Predictions and Intervals

PREDICTIONS AND INTERVALS ANALYSIS

DATASET OVERVIEW
Dataset shape: (506, 14)
Column names: ['crim', 'zn', 'indus', 'chas', 'nox', 'rm', 'age', 'dis', 'rad',
'tax', 'ptratio', 'b', 'lstat', 'medv']

Using 'medv' as target variable
Using 'lstat' as predictor variable (lstat)

=== SIMPLE LINEAR REGRESSION MODEL ===
Model: medv ~ lstat
R-squared: 0.5441
Regression equation: medv = 34.5538 + -0.9500 × lstat

=== 5.1: PREDICTION FOR LSTAT = 10% ===
CALCULATION:
ŷ = β_0 + β_1 × X
ŷ = 34.5538 + -0.9500 × 10.0
ŷ = 25.0533

Predicted median home value for lstat = 10%: $25.05k

=== 5.2: 95% CONFIDENCE INTERVAL FOR MEAN RESPONSE ===
CALCULATION DETAILS:

```
• Predicted value: 25.0533
• Standard error of mean: 0.2948
• t-critical (α=0.05, df=504.0): 1.9647
• Margin of error: 0.5792

95% CONFIDENCE INTERVAL: [24.4741, 25.
  ↪6326]
In dollars: [$24.47k, $25.63k]

INTERPRETATION:
We are 95% confident that the mean␣
  ↪median home value for all␣
  ↪neighborhoods
with lstat = 10% is between $24.47k␣
  ↪and $25.63k.


=== 5.3: 95% PREDICTION INTERVAL FOR␣
  ↪INDIVIDUAL RESPONSE ===
CALCULATION DETAILS:
• Predicted value: 25.0533
• Standard error of prediction: 6.4803
• t-critical (α=0.05, df=504.0): 1.9647
• Margin of error: 12.7316

95% PREDICTION INTERVAL: [12.3217, 37.
  ↪7850]
In dollars: [$12.32k, $37.78k]

INTERVAL COMPARISON:
• Confidence interval width: 1.1584
• Prediction interval width: 25.4633
• Prediction interval is 21.98x wider␣
  ↪than confidence interval
```

5.4 Explain the difference between a confidence interval and a prediction interval in practical terms. When would you use each type?

todo

```
=== 5.4: CONFIDENCE VS PREDICTION␣
  ↪INTERVALS ===

CONCEPTUAL DIFFERENCES:

CONFIDENCE INTERVAL:
• Estimates uncertainty about the MEAN␣
  ↪response for a given X value
• Answers: 'What is the average Y for␣
  ↪all observations with this X?'
• Accounts for uncertainty in␣
  ↪estimating the population mean
• Gets narrower as sample size␣
  ↪increases
• Narrower interval (less uncertainty)

PREDICTION INTERVAL:
• Estimates uncertainty about an␣
  ↪INDIVIDUAL response for a given X␣
  ↪value
• Answers: 'What might Y be for a␣
  ↪single new observation with this X?'
• Accounts for both estimation␣
  ↪uncertainty AND individual variation
• Includes natural scatter around the␣
  ↪regression line
• Wider interval (more uncertainty)

WHEN TO USE EACH:
--------------------

USE CONFIDENCE INTERVAL when:
• Estimating average outcomes for␣
  ↪policy/planning
• Comparing mean responses between␣
  ↪groups
• Making statements about population␣
  ↪parameters
• Example: 'What's the average home␣
  ↪value in 10% lstat neighborhoods?'

USE PREDICTION INTERVAL when:
• Predicting outcomes for specific␣
  ↪individuals/cases
• Setting bounds for individual␣
  ↪forecasts
• Risk assessment for single␣
  ↪observations
• Example: 'What might this specific␣
  ↪house be worth?'
```

5.5 For lstat values of 5%, 15%, and 25%, calculate point predictions and comment on how the relationship changes across different levels of the predictor variable

todo

```
=== 5.5: PREDICTIONS AT MULTIPLE LSTAT VALUES ===

POINT PREDICTIONS:
lstat = 5%:
  → Predicted value: $29.80k
  → 95% CI: [$29.01k, $30.60k]
  → 95% PI: [$16.63k, $42.98k]

lstat = 10%:
  → Predicted value: $25.05k
  → 95% CI: [$24.47k, $25.63k]
  → 95% PI: [$12.32k, $37.78k]

lstat = 15%:
  → Predicted value: $20.30k
  → 95% CI: [$19.73k, $20.87k]
  → 95% PI: [$7.58k, $33.02k]

lstat = 25%:
  → Predicted value: $10.80k
  → 95% CI: [$9.72k, $11.89k]
  → 95% PI: [$-3.15k, $24.75k]

RELATIONSHIP ANALYSIS:

Model slope (β1): -0.9500
Interpretation: For each 1% increase in lstat, median home value
decreases by $0.95k on average

CHANGES BETWEEN LSTAT LEVELS:
• 5.0% → 10.0%: Change = $-4.75k
  Rate: $-0.95k per 1% lstat increase
• 10.0% → 15.0%: Change = $-4.75k
  Rate: $-0.95k per 1% lstat increase
• 15.0% → 25.0%: Change = $-9.50k
  Rate: $-0.95k per 1% lstat increase

COMMENTS ON RELATIONSHIP:
• The relationship shows moderate negative association
• Linear relationship assumed constant across all lstat levels
• Higher lstat (more lower status population) associated with lower home values
```

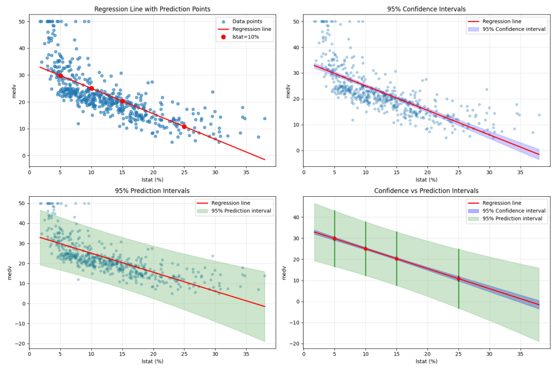


```
=== PREDICTIONS SUMMARY TABLE ===

DETAILED PREDICTIONS TABLE:
 lstat  prediction  ci_lower  ci_upper  pi_lower  pi_upper  ci_width  pi_width
width_ratio
     5      29.804    29.007    30.600    16.627    42.980     1.592    26.353
16.550
    10      25.053    24.474    25.633    12.322    37.785     1.158    25.463
21.981
    15      20.303    19.732    20.875     7.585    33.021     1.143    25.436
22.254
    25      10.803     9.717    11.888    -3.148    24.754     2.170    27.902
12.856

KEY INSIGHTS:
• As lstat increases, predicted home values decrease
```

```
• Prediction intervals are
  consistently 18.4x wider than
  confidence intervals
• The linear relationship appears
  moderate (R^2 = 0.544)

=== MODEL ASSUMPTIONS REMINDER ===
For these intervals to be valid,
  ensure:
• Linear relationship between variables
• Independence of residuals
• Homoscedasticity (constant variance)
• Normality of residuals
• No influential outliers
```